

Obliczenia superkomputerowe - HPC (Info II)

18/19

Kokpit ► Moje kursy ► HPC.INFOII.18/19 ► CUDA ► CUDA assignment 2019

NAWIGACJA

Kokpit

Strona główna

Strony

Moje kursy

TA.ZOB.19/20

HPC.INFOII.18/19

Uczestnicy

Odznaki

Kompetencje

Oceny

Sekcja

ogólna

Introduction to HPC

OpenMP

CUDA

lecture notes: 02 - CUDA

Lab 3 - CUDA basics

Lab 4 - CUDA memory types, synchronization

Lab 5 - advanced features of CUDA

CUDA assignment 2019

Massively Parallel A* Search on a GPU

1. Introduction

A* (or AStar) is a graph search algorithm widely used in artificial intelligence. While the traditional version of this algorithm is purely sequential, there have been several parallel versions over the years. One such variant is the approach described by *Zhou* and *Zeng* in their paper *Massively Parallel A* Search on a GPU*. It's a version of A* designed for GPU execution. Your task in this assignment is to implement that algorithm with CUDA.

2. Details

The algorithm is described in this paper: *Massively Parallel A* Search on a GPU*.

The paper describes the search algorithm itself and two node duplication detection schemes: *Parallel Cuckoo Hashing* and *Parallel Hashing with Replacement*. Duplication detection is needed so the A* algorithm doesn't expand already visited nodes. In your solutions you need to focus only on *Parallel Hashing with Replacement* - it's a bit easier to implement. You can find the pseudocode for that duplication detection scheme in the appendix (or below in Figure 2).

- 19. CUDA assignment 2019
- CUDA debugging, profiling and optimisation
- CUDA assignment 2019
- CUDA assignment - solution form
- Parallel algorithms in the PRAM model
- Quantitative Efficiency Models
- MPI
- Modern Supercomputers: architecture and system sof...
- Algorithms in the Latency-Bandwidth model
- TBB: Threading Building Blocks
- Programming HPC machines
- Fundamental Distributed Algorithms
- Scheduling
- Summary and perspectives

```

1: procedure GA*( $s, t, k$ )
     $\triangleright$  find the shortest path from  $s$  to  $t$  with  $k$  queues
2:   Let  $\{Q_i\}_{i=1}^k$  be the priority queues of the open list
3:   Let  $H$  be the closed list
4:   PUSH( $Q_1, s$ )
5:    $m \leftarrow \text{nil}$   $\triangleright m$  stores the best target state
6:   while  $Q$  is not empty do
7:     Let  $S$  be an empty list
8:     for  $i \leftarrow 1$  to  $k$  in parallel do
9:       if  $Q_i$  is empty then
10:        continue
11:       end if
12:        $q_i \leftarrow \text{EXTRACT}(Q_i)$ 
13:       if  $q_i.\text{node} = t$  then
14:         if  $m = \text{nil}$  or  $f(q_i) < f(m)$  then
15:            $m \leftarrow q_i$ 
16:         end if
17:         continue
18:       end if
19:        $S \leftarrow S + \text{EXPAND}(q_i)$ 
20:     end for
21:     if  $m \neq \text{nil}$  and  $f(m) \leq \min_{q \in Q} f(q)$  then
22:       return the path generated from  $m$ 
23:     end if
24:      $T \leftarrow S$ 
25:     for  $s' \in S$  in parallel do
26:       if  $s'.\text{node} \in H$  and  $H[s'.\text{node}].g < s'.g$  then
27:         remove  $s'$  from  $T$ 
28:       end if
29:     end for
30:     for  $t' \in T$  in parallel do
31:        $t'.f \leftarrow f(t')$ 
32:       Push  $t'$  to one of priority queues
33:        $H[t'.\text{node}] \leftarrow t'$ 
34:     end for
35:   end while
36: end procedure

```

scores

MRJP.INFO.II.1

8/19

PSZI.INFO.III.1

7/18

JiPP.INFO.III.17

/18

ZPP.INFO.III.17

/18

IO.INFO.II.16/1

7

JNPI.INFO.II.16

/17

BD.INFO.II.16/1

7

**ADMINISTRACJA**

Administracja
kursem

```

1: procedure HASH-WITH-REPLACEMENT-DEDUPLICATE( $H, T$ )
2:    $T' \leftarrow T$ 
3:   for  $i \leftarrow 0$  to  $|T|$  in parallel do
4:      $z \leftarrow 0$ 
5:     for  $j \leftarrow 0$  to  $d - 1$  do
6:       if  $H[h_j(T[i])] \in \{T[i], \text{nil}\}$  then
7:          $z \leftarrow j$ 
8:         break
9:       end if
10:    end for
11:     $t \leftarrow T[i]$ 
12:    ATOMIC-SWAP( $t, H[h_z(T[i])]$ )
13:    if  $t = T[i]$  then
14:      remove  $T[i]$  from  $T'$ 
15:      continue
16:    end if
17:    for  $j \leftarrow 0$  to  $d - 1$  do
18:      if  $j \neq z$  and  $H[h_j(T[i])] = T[i]$  then
19:        remove  $T[i]$  from  $T'$ 
20:        break
21:      end if
22:    end for
23:  end for
24:  return  $T'$ 
25: end procedure

```

Figure 1. Pseudocode for the parallel A* algorithm**Figure 2.**

Pseudocode for Parallel Hashing with Replacement

The paper presents 3 possible applications of A*: sliding puzzles, pathfinding and protein design. In this assignment we'll focus on the first two.

a) Sliding puzzles

A sliding puzzle is a task of finding the smallest number of moves from one configuration of cells to another in an $N \times N$ grid, where each move slides one cell to an empty space in the grid. Here we'll use $N = 5$.

1	2	3	9	4
6	14	19	23	5
16	22	18		8
21	12	7	15	13
17	11	20	24	10

We need to use a heuristic function that approximates a distance to the target configuration. The *Zhou* paper uses the *disjoint pattern database* method, but we'll use a simple *Manhattan* distance instead. Let's look at an example for a 3x3 grid:

2 1 3	1 2 3	1 2 3 4 5 6 7 8
5 4 0	4 5 6	-----
6 7 8	7 8 0	1+1+0+1+1+3+1+1 = 9

For each cell we count the number of moves necessary to get to their target positions (or simply $|x_1 - x_2| + |y_1 - y_2|$, where x_i, y_i are cell coordinates) and sum them together. The resulting number is an approximate distance between two configurations.

b) Pathfinding

In this task we search for a shortest path between two points on a grid. Each cell is connected with up to 8 neighbours and connections are weighted. One example of a task like this is moving a unit in the *Civilization* video game. The map in that game is comprised of square tiles. Each of them has a specific terrain type which has an effect on movement speed. For instance, moving through mountains is more difficult than over plains, so the movement cost is higher in the former case.




Figure 4. Map view in the original *Civilization*

In this assignment, all weights are 1 unless they're overridden in the input file. Similarly, it's possible to set obstacles in the input file - nodes don't form connections to cells denoted as such. Cell at position $(0, 0)$ is in the lower left corner of the grid. Weight values are integers, ranging from 1 to 5. The maximum grid size is 10000 x 10000. Cell coordinates are formatted as (x, y) .

For the heuristic function, we again use the *Manhattan* distance:

$$|x_1 - x_2| + |y_1 - y_2|.$$

3. Requirements

Write your kernel code in CUDA C. The host code can be either in C or C++. The specifics of the device code are up to you - choose appropriate memory types, synchronisation techniques, use streaming or not, etc. You can optimise for Titan X GPUs on *bruce*. Use one GPU - you don't have to utilise a multi-GPU setup.

IMPORTANT: You can't use third-party libraries in your kernels - all code that's executed on the GPU needs to be implemented by you. You can however use external libraries in your host code to handle I/O tasks, prepare input data, etc.

You can organise your source code in any way you like, but it has to be properly divided into modules. Your Makefile should produce a single executable named `astar_gpu`.

That executable must handle the following command line arguments:

```
--version sliding | pathfinding - selects either sliding puzzle or pathfinding
--input-data PATH - file containing graph definition
--output-data PATH - file containing program execution time in milliseconds
(excluding I/O operations) and search results
```

You should describe your solution in a detailed report - focus on explaining your implementation and optimisation techniques. Include screenshots from NVIDIA Visual Profiler that show GPU utilisation during program execution.

Send your solution as a zip file named `ab123456_CUDA.zip`, where `ab123456` is your login (use this form). It should have the following contents:

```
ab123456_CUDA/
  source/
    src_file1.cu
    src_file2.cu
    ...
  Makefile
  report/
    ab123456_report.pdf
```

a) Sliding puzzle

```
./astar_gpu --version sliding --input-data /home/user/input.txt --output-data /home/user/result.txt
```

```
1, 2, _, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23,
24
4, 9, 8, 12, 14, _, 18, 13, 11, 19, 21, 23, 20, 24, 15, 1, 3, 7, 16, 5, 6, 2, 10, 17,
22
```

The first line is the start configuration and the second line is the target configuration.

The output file should look like this:

```
1324
1, 2, _, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23,
24
1, 2, 3, _, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23,
24
...
4, 9, 8, 12, 14, 18, _, 13, 11, 19, 21, 23, 20, 24, 15, 1, 3, 7, 16, 5, 6, 2, 10, 17,
22
4, 9, 8, 12, 14, _, 18, 13, 11, 19, 21, 23, 20, 24, 15, 1, 3, 7, 16, 5, 6, 2, 10, 17,
22
```

The first line is the execution time in milliseconds. The following lines are the shortest path found by the A* algorithm.

b) Pathfinding

```
./astar_gpu --version pathfinding --input-data /home/user/input.txt --output-data
/home/user/result.txt
```

For pathfinding, the input file must have the following format:

```
10, 10
0, 0
9, 9
2
1, 1
2, 3
4
2, 2, 3
4, 3, 2
1, 4, 5
3, 4, 2
```

The first line is the grid size. The second line is the start position, followed by the end position in the third line. The fourth line is the number of obstacle positions, o . Then there are o lines with obstacle coordinates. This is followed by a line containing the number of cells which have connections with weights different than 1, w . Then there are w lines, each storing the cell coordinates and a weight value.

The output file should look like this:

```
2345
0,0
0,1
...
8,9
9,9
```

The first line is the execution time in milliseconds. The following lines are the shortest path found by the A* algorithm.

c) Path doesn't exist

If the algorithm can't find a solution (either in the sliding puzzle or pathfinding task), the program should finish properly and create an output file that contains only the execution time in milliseconds.

4, Example data

a) Sliding puzzle

Input: puzzle_input.txt

Output: puzzle_output.txt

b) Pathfinding

Input: pathfinding_input.txt

Output: pathfinding_output.txt

5. Frequently Asked Questions

1) What's the submission deadline?

28.04.2019 23:55

2) What is the upper limit for the grid size in the pathfinding problem?

10000 x 10000 cells, for more details see section **2b**.

Edit history:

29.03.2019: Initial version

15.04.2019: Added limits for graph weights and grid size in the pathfinding problem. Added FAQ.

19.04.2019: Added example input and output data. Added requirements for the case when the path doesn't exist.

Ostatnia modyfikacja: piątek, 19 kwiecień 2019, 21:09

◀ Lab 6 - CUDA
debugging, profiling
and optimisation

Przejdź do...

CUDA assignment -
solution form ▶

Jesteś zalogowany(a) jako Jakub Bujak (Wyloguj)

HPC.INFOII.18/19

Podsumowanie zasad przechowywania danych

Pobierz aplikację mobilną

Moodle, wersja 3.5.7+ (Build: 20190823) | moodle@mimuw.edu.pl